


NA62 online pc farm design and implementation

Jonas Kunze


Universität Mainz

05.12.2011



E226
NA62
8748

- 1 Motivation
 - The physics
- 2 The experiment
 - Subdetectors
- 3 DAQ and Trigger
 - Data acquisition
 - Three level online trigger
- 4 Test system
 - Components bought for testing
- 5 Merging L1 and L2
 - New farm design
 - Event building @ L1
- 6 Outlook and conclusion



$K^+ \rightarrow \pi^+ \nu \bar{\nu} \Leftrightarrow V_{td}$ of CKM matrix

SM branching ratio: $(8.5 \pm 0.7) \cdot 10^{-11}$

Measurement from BNL (Brookhaven)

The only experimental data so far (7 candidates) E787 and E949:

$$BR(K^+ \rightarrow \pi^+ \nu \bar{\nu}) = (17.3^{+11.5}_{-10.5}) \cdot 10^{-11}$$

NA62 aims $\sigma < 10\%$ with 100 events

$\approx 10^{13}$ K^+ decays required

High efficiency needed \Rightarrow high data rate

Signal-to-noise ratio of 1/10 planned with an event rate of 10 MHz


Background channels to be suppressed by the detector

$K^+ \rightarrow \mu^+ \nu$	(63%) $\rightarrow \mu$ veto
$K^+ \rightarrow \pi^+ \pi^0$	(21%) $\rightarrow \gamma$ veto
$K^+ \rightarrow \pi^+ \pi^+ \pi^-$	(6%) \rightarrow charged particle veto
$K^+ \rightarrow \pi^0 e^+ \nu$	(5%) $\rightarrow \gamma$ veto
$K^+ \rightarrow \pi^0 \mu^+ \nu$	(3%) $\rightarrow \gamma, \mu$ veto
$K^+ \rightarrow \pi^+ \pi^0 \pi^0$	(2%) $\rightarrow \gamma$ veto




Detector dominated by vetos

Experiment overview



CEDAR




Differential Cerenkov counter

K^+ identification

Mirrors collect K^+ -specific Cerenkov light


GTK



Gigatracker (second Achromat)

Timing, momentum and angle of all particles using silicon pixels
750 MHz event rate!!!


Straws



Spectrometer

Tracking of charged particles within a magnetic field

Photon vetos part 1: LAV, IRC, SAC




Photon vetos

Huge range of angle to be covered

Scintillators at different positions (lead-glass reused from OPAL)

Photon veto part 2: LKr




Liquid Krypton calorimeter

Photon veto for central angles range

Muon veto

Reused from NA48 \Rightarrow can be used as calorimeter for analysis of additional rare decays

Rich




Ring Imaging Cherenkov

Separate π^+ from μ^+ between 15 and 35 GeV/c (L0 trigger)

NA62
P226

MUV



Muon veto

Scintillator-iron-sandwich mainly produced in Mainz


NA62
87M

Data rates

10 MHz unsteady K^+ rate

Detector	Event size [B]	Data rate [GBps]
CEDAR	216	2.16
GTK	2250	22.50
CHANTI	192	1.92
LAV	160	1.60
STRAW	768	7.68
RICH	160	1.60
CHOD	$\ll 1000$	$\ll 10$
MUV	768	7.68
IRC & SAC	576	5.76
LKR	222 k	2220
Sum	≈ 227 kB	≈ 2.3 TBps

The TEL-62 board



Board with Altera Stratix III FPGAs

- Based on TELL-1 from LHCb
- 1 Credit card PC
- 4+1 Gigabit ethernet interfaces
- 2 or 4 GB DDR2 memory
- Up to 4 TDC daughter cards


We will have about 500 Tel-62 and similar boards at NA62

These boards are used for DAQ and for L0 triggering

DAQ and Trigger system

Three levels to filter data

Data transmission via ordinary 10 gigabit ethernet:



NA62
P226

The Triggers

Original concept:

Level 0 (hardware)

≈500 FPGA boards (mainly Tel-62) downstairs at the experiment

Level 1 (software, on subdetector level)

≈10 Subdetector specific PCs (8 cores with GPUs) upstairs (up to 100m fibre for 10G ethernet)

Level 2 (software, whole detector)

≈30 PCs (8 cores) that merge the subdetector data to one event (eventbuilding) and process a third trigger decision

The Triggers - assignments

Level 0

Hit in the Rich, no γ and no μ

Level 1

Every subdetector decides separately

Mainly data integrity checks

Ring fitting of the RICH via GPUs

Track fitting for straws using GPUs planned

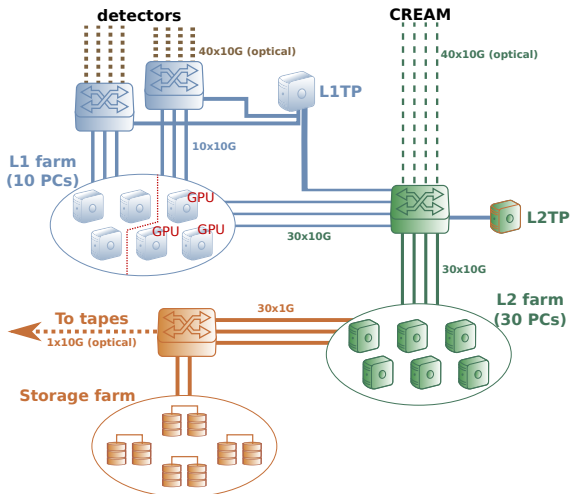
Level 2

Eventbuilding (merging of all subdetector data) and reconstruction



First topology proposal

Original concept:



Special hardware purchased for performance tests

High performance 12-core PC Dell R710

2*6 Cores @ 2.93GHz (3.33GHz turbo)

24GB memory @ 1333MHz

Low power 12-core PC Dell R710

2*6 Cores @ 2.26GHz (2.8GHz turbo)

24GB memory @ 1333MHz

48 x 1G-port switch Dell PowerConnect 6248

Two SFP+ (10 Gbps) modules (4 ports total)

Test farm @ Mainz



2 big 12 core machines for farm simulation

Connected via 10Gbps SFP+ to the switch
and to each other

11 L0 simulator PCs (1-4 cores)

Diskless

21 Cores total

21 * 1Gbps links connected to the switch
(up to 4Gbps per PC)

Burst time and duty-cycle

Only 3-9 sec. burst and long break

Duty cycle: $T_{Burst}/T_{Break} \approx 0.3$




Abb.: NA48 approach

Reuse L1 PCs

My proposal to use resources more efficiently

Reuse L1 PCs during burst break for L2 computation by combining L1 and L2 to one farm



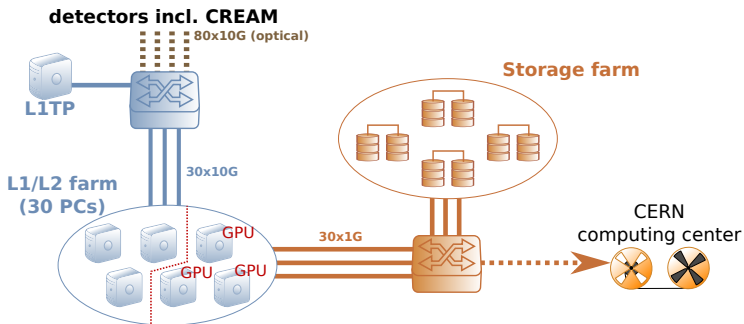


Abb.: New proposal

Don't separate L1 and L2!




Combine L1 and L2 to one farm



We safe about 80k

- No L1 PCs anymore
- Less switches, less network cards

New bottleneck




Risk of overloading the farm (L1 and L2 concurrently)

Possibilities to solve this problem:

- Implement a very intelligent load balancing algorithm
- Wait with L2 until end of Burst (L1 has finished)

New proposal

Event building @ L1




Every subdetector sends data of an event to one single PC

- + No broadcast of a L1 decision needed anymore (no L1TP)
- + Easier to implement load balancing (self-sustaining PCs)
- Every farm PC must serve every subdetector \Rightarrow needs GPUs

UDP Storm

Congestion at one single farm PC



Bundle events for better performance

Smallest Event: 160 B

Jumbo ethernet frame: 9 kB

55 Events in one bundle

$\Rightarrow 55 \cdot 5kB = 275kB$ total per bundle

Round Robin tables at L0


Send events 1-55 to PC A

Send events 56-110 to PC B

...

Packet loss

22 L0 boards simulated



Packet loss

Under construction

High packet loss

Packets are rejected by Kernel at receiver side (Not at switch or network card)!

I'm investigating if we can use standard drivers or have to implement own ones

Packets get lost in software, not hardware!

⇒ UDP storms don't cause problems

Events can be merged at L1 which makes load balancing much easier (no L1TP needed)

No expensive core switch needed

One expensive core switch was originally planned

- + Many ports (high density)
- + Fast backplane (any to any @ 10Gbps bidirectional)
- High cost (about 180k)


Tree topology with cheap 24 and 48 port switches

- Slow backplanes (any to any @ 10Gbps only unidirectional)
- Bottleneck at switch-to-switch connection
- + Low cost (about 80k)

Only one way transmission with about 8Gbps per link

It is feasible to use slower switches

Tree topology (Hexapus)



NA62
P226
NA48

Current investigations and outlook

Implementation of a high performance L1/L2 farm software

- Data transmission L0 → farm and L2 → tapes
- Interface for the actual trigger algorithms
- Monitoring and administration system with a complex web interface

Further performance tests


- Do we need special network drivers to reduce packet loss?
- How many PCs will be needed?

Conclusion


- High energy and high precision \Rightarrow a lot of data
- Using ordinary ethernet saves money and time and gives you the ability to quickly switch between different approaches
- Unsteady data production allows new approaches
 - Considering trigger levels as logical object, not as real farms saves a lot of money
- The new farm design allows us to have a central software architecture which is much easier to implement \Rightarrow Now Mainz will manage the whole chain from L0 to persistent memory and install one central farm instead of only L2

Thank you!


2097152B memory - TCP



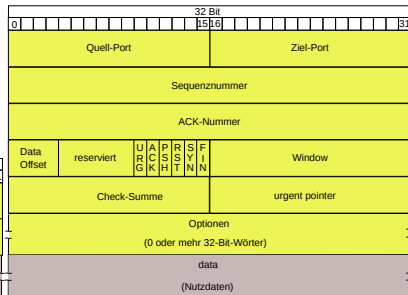
2097152B memory - UDP



TCP vs. UDP




UDP-Header

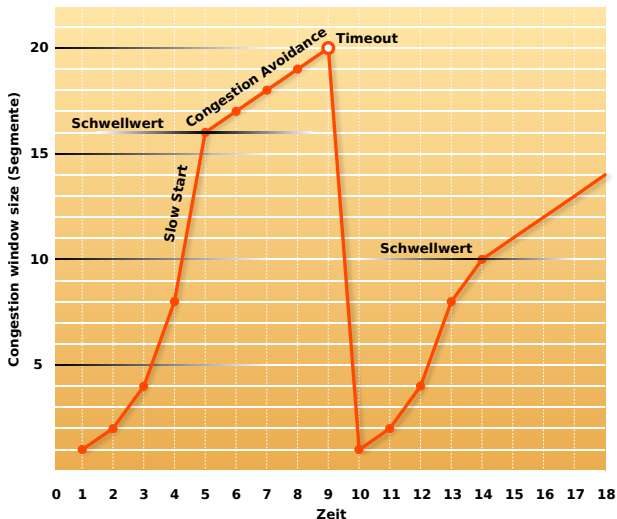


TCP-Header

TCP: Reliability and flow control



Congestion Avoidance



Performance tests



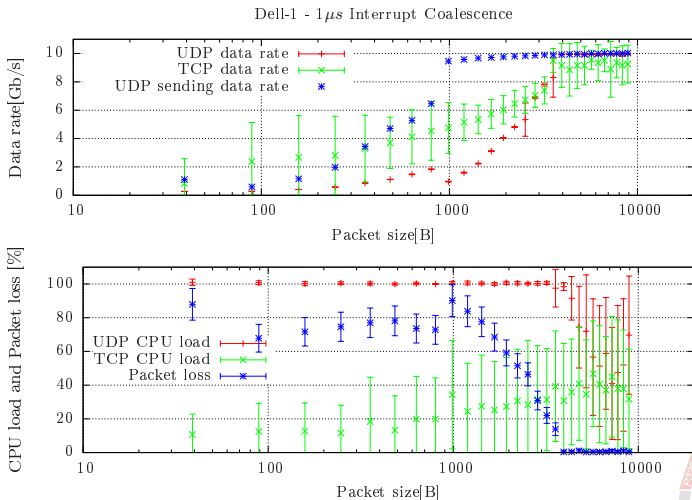
TCP optimizes the usage of network resources

But what does this cost?

Intuitively one would guess: Higher CPU usage and longer latencies.

but...

Data rate and CPU usage



TCP Offload Engine


Network cards and drivers are optimized for TCP

- Checksums and fragmentation calculated on hardware
- Some drivers ignore interrupt coalescence of $0\mu s$



Using TCP reduces CPU usage \Rightarrow more space for computation

Timing



Results

- TCP reduces CPU usage \Rightarrow more space for computation
- TCP has flow control and congestion avoidance
- TCP is reliable

TCP in FPGAs

TCP means high payload in hardware!

\Rightarrow TCP can only be used for PC to PC communication at NA62!

TCP/UDP vs. basic IP

Using standard interrupt/kernel based socket programming...

- is optimized for TCP (drivers)
- + is easy and many libraries can be used (e.g. boost::asio)
- induces high latency ($\approx 30 - 150\mu s$)
- induces high packet loss??!

Programming own Kernel modules...

- is hard stuff (only few small libraries)
- is bound to hardware
- + highest performance possible